

DroneDetect V2 - CNN Classification (Google Colab)

Train CNN classifiers on spectrogram features:

- VGG16 (frozen features + trainable FC)
- ResNet50 (frozen features + trainable FC)
- File-level stratified split to prevent data leakage
- Side-by-side performance comparison

1. Mount Google Drive

```
In [22]: from google.colab import drive  
  
drive.mount('/content/drive')
```

```
Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).
```

2. Imports

```
In [23]: !pip install -U kaleido==0.2.1  
  
Requirement already satisfied: kaleido==0.2.1 in /usr/local/lib/python3.12/dist-packages (0.2.1)
```

```
In [24]: import numpy as np  
import matplotlib.pyplot as plt  
import seaborn as sns  
import plotly.express as px  
import plotly.graph_objects as go  
from plotly.subplots import make_subplots  
import torch  
import torch.nn as nn  
import torch.optim as optim
```

```

from torch.utils.data import Dataset, DataLoader
from sklearn.model_selection import StratifiedGroupKFold
from sklearn.metrics import confusion_matrix, classification_report, accuracy_score, f1_score
from tqdm import tqdm
import os
import gc
import psutil
import re
from pathlib import Path

import torchvision.models as tv_models

plt.style.use('seaborn-v0_8-darkgrid')
%matplotlib inline

device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
print(f"Using device: {device}")

# Figure saving configuration
NOTEBOOK_NAME = "training_cnn_COLAB"
FIGURES_DIR = Path("figures") / NOTEBOOK_NAME

def save_figure(fig) -> None:
    """Save plotly figure to PNG file using the figure's title as filename."""
    FIGURES_DIR.mkdir(parents=True, exist_ok=True)
    title = fig.layout.title.text if fig.layout.title.text else "untitled"
    filename = re.sub(r'^\w\s-]', '', title).strip()
    filename = re.sub(r'[\s-]+', '_', filename)
    filepath = FIGURES_DIR / f"{filename}.png"
    try:
        fig.write_image(str(filepath), width=1200, height=800)
        print(f"Saved: {filepath}")
    except Exception as e:
        print(f"Warning: Could not save figure (kaleido required): {e}")

```

Using device: cuda

3. Configuration

```
In [25]: CONFIG = {
    # Paths
    'features_path': 'drive/MyDrive/DroneDetect_V2/output/features/spectrogram_features.npz',
    'models_dir': 'drive/MyDrive/DroneDetect_V2/output/models/',
    'test_data_dir': 'drive/MyDrive/DroneDetect_V2/output/sample/test_data/',

    # Split parameters
    'test_size': 0.2,
    'random_state': 42,

    # Training parameters
    'batch_size': 128,
    'epochs': 10,
    'learning_rate': 0.01,

    # Device
    'device': device
}

print(f"Configuration: {CONFIG}")

Configuration: {'features_path': 'drive/MyDrive/DroneDetect_V2/output/features/spectrogram_features.npz', 'models_dir': 'drive/MyDrive/DroneDetect_V2/output/models/', 'test_data_dir': 'drive/MyDrive/DroneDetect_V2/output/sample/test_data/', 'test_size': 0.2, 'random_state': 42, 'batch_size': 128, 'epochs': 10, 'learning_rate': 0.01, 'device': device(type='cuda')}
```

4. Model Definitions

```
In [26]: class VGG16FC(nn.Module):
    """VGG16 with frozen features and trainable classifier.

    This model uses a pre-trained VGG16 backbone with weights frozen,
    replacing the classifier with a new fully connected layer for the
    specific number of classes.

    Attributes:
        features (nn.Sequential): The feature extractor part of VGG16.
        classifier (nn.Linear): The trainable classification layer.
    """

```

```
def __init__(self, num_classes: int):
    """Initializes VGG16FC.

    Args:
        num_classes (int): The number of output classes for classification.
    """
    super().__init__()
    # Removed from_array logic as inputs are strictly RGB

    vgg = tv_models.vgg16(weights='IMAGENET1K_V1')
    self.features = nn.Sequential(*list(vgg.children())[:-1])

    for param in self.features.parameters():
        param.requires_grad = False

    self.classifier = nn.Linear(25088, num_classes)

    def forward(self, x: torch.Tensor) -> torch.Tensor:
        """Forward pass of the model.

        Handles input permutation from NHWC to NCHW format if necessary.

        - NHWC: (Batch, Height, Width, Channels) - Common in NumPy/OpenCV.
        - NCHW: (Batch, Channels, Height, Width) - Standard for PyTorch.

        Args:
            x (torch.Tensor): Input tensor. Shape can be (N, H, W, C) or (N, C, H, W).
        Returns:
            torch.Tensor: Model logits.
        """
        # Handle RGB inputs: NHWC -> NCHW
        if x.dim() == 4 and x.shape[-1] == 3:
            x = x.permute(0, 3, 1, 2)

        x = self.features(x)
        # Use flatten(1) to handle non-contiguous tensors and flatten starting from batch dim
        x = x.flatten(1)
        return self.classifier(x)

    def reset_weights(self):
```

```
    """Resets the weights of the classifier layer."""
    self.classifier.reset_parameters()

class ResNet50FC(nn.Module):
    """ResNet50 with frozen features and trainable classifier.

    This model uses a pre-trained ResNet50 backbone with weights frozen.
    The final fully connected layer and adaptive pooling are removed and
    replaced with a new linear classifier.

    Attributes:
        features (nn.Sequential): The feature extractor part of ResNet50.
        classifier (nn.Linear): The trainable classification layer.
    """

    def __init__(self, num_classes: int):
        """Initializes ResNet50FC.

        Args:
            num_classes (int): The number of output classes for classification.
        """
        super().__init__()

        resnet = tv_models.resnet50(weights='IMAGENET1K_V1')
        # Remove FC and adaptive pooling
        self.features = nn.Sequential(*list(resnet.children())[:-2])

        for param in self.features.parameters():
            param.requires_grad = False

        # Output of ResNet50 before pooling: 2048 x 7 x 7 = 100352
        self.classifier = nn.Linear(100352, num_classes)

    def forward(self, x: torch.Tensor) -> torch.Tensor:
        """Forward pass of the model.

        Handles input permutation from NHWC to NCHW format if necessary.

        - NHWC: (Batch, Height, Width, Channels) - Common in NumPy/OpenCV.
        - NCHW: (Batch, Channels, Height, Width) - Standard for PyTorch.
    
```

```

Args:
    x (torch.Tensor): Input tensor. Shape can be (N, H, W, C) or (N, C, H, W).

Returns:
    torch.Tensor: Model logits.
"""
if x.dim() == 4 and x.shape[-1] == 3:
    x = x.permute(0, 3, 1, 2)

x = self.features(x)
# Use flatten(1) to handle non-contiguous tensors and flatten starting from batch dim
x = x.flatten(1)
return self.classifier(x)

def reset_weights(self):
    """Resets the weights of the classifier layer."""
    self.classifier.reset_parameters()

print("Model classes defined (VGG16FC, ResNet50FC)")

```

Model classes defined (VGG16FC, ResNet50FC)

4b. Memory-Efficient Dataset

Problem: Spectrogram NPZ file is 11.73 GB (19478 samples x 224x224x3 x 4 bytes). Using fancy indexing `X_train = X[train_idx]` on a memory-mapped array forces NumPy to load the entire subset into RAM, causing OOM on Colab.

Solution: Custom `RGBMemmapDataset` that stores indices and loads one sample at a time via `__getitem__`. The `.copy()` call breaks the memmap view and loads only the requested sample.

Memory savings: 11.73 GB upfront load → ~77 MB per batch (128 samples) + 2 GB PyTorch overhead. Total RAM usage: < 5 GB vs 11.73 GB.

Trade-off: Disk I/O overhead (~10-20s per epoch), but prevents OOM crashes.

```
In [27]: class RGBMemmapDataset(Dataset):
    """
    Memory-efficient dataset for memmap RGB spectrograms.

```

Loads one sample at a time from disk instead of loading entire array into RAM.

Memory optimization: Avoids fancy indexing on memmap which would load entire subset into RAM (11.73 GB). Instead, stores indices and loads samples individually via `__getitem__`.

Memory usage: ~77 MB per batch (128 samples) vs 11.73 GB upfront.

Parameters

```
memmap_array : np.memmap
    Memory-mapped array from np.load(..., mmap_mode='r')
indices : np.ndarray
    Indices for this split (train or test)
labels : np.ndarray
    Labels for samples
"""

def __init__(self, memmap_array, indices, labels):
    self.memmap = memmap_array
    self.indices = indices
    self.labels = labels

def __len__(self):
    return len(self.indices)

def __getitem__(self, idx):
    # Map dataset index to original array index
    actual_idx = self.indices[idx]
    # CRITICAL: .copy() breaks memmap view and Loads single sample
    rgb = self.memmap[actual_idx].copy()  # (224, 224, 3)

    # Convert to tensor
    x = torch.from_numpy(rgb).float()
    y_label = torch.tensor(self.labels[idx]).long()

    return x, y_label

print("Memory-efficient RGBMemmapDataset class defined")
```

Memory-efficient RGBMemmapDataset class defined

5. File-Level Stratified Split Function

```
In [28]: def get_stratified_file_split(X, y, file_ids, test_size=0.2, random_state=42):
    """
    Split data at FILE level to prevent data leakage.

    Segments from the same .dat file (~100 segments) will never appear
    in both train and test sets.

    Parameters
    -----
    X : array-like
        Features (n_samples, ...)
    y : array-like
        Labels for stratification (n_samples,)
    file_ids : array-like
        Source file ID for each sample (n_samples,)
    test_size : float
        Approximate test set proportion (actual may vary due to file grouping)
    random_state : int
        Random seed for reproducibility

    Returns
    -----
    train_idx, test_idx : arrays
        Indices for train/test split
    """
    n_splits = int(1 / test_size) # e.g., test_size=0.2 -> 5 splits -> 1 fold = 20%

    sgkf = StratifiedGroupKFold(n_splits=n_splits, shuffle=True, random_state=random_state)

    # Take first fold as train/test split
    train_idx, test_idx = next(sgkf.split(X, y, groups=file_ids))

    # Verify no file leakage
    train_files = set(file_ids[train_idx])
    test_files = set(file_ids[test_idx])
```

```

    assert len(train_files & test_files) == 0, "Data leakage detected: files in both splits"

    return train_idx, test_idx

print("Split function defined")

```

Split function defined

6. Streaming Dataset Definition

```
In [29]: # Memory optimization: use memory mapping to avoid loading full array
data = np.load(CONFIG['features_path'], mmap_mode='r', allow_pickle=True)

X_memmap = data['X'] # Shape: (N, 224, 224, 3) - memory mapped, not loaded
y_drone = data['y_drone'][::] # Load labels (small, ~76 KB)
y_interference = data['y_interference'][::]
y_state = data['y_state'][::]
file_ids = data['file_ids'][::] # Load file IDs (small, ~76 KB)
drone_classes = data['drone_classes']
interference_classes = data['interference_classes']
state_classes = data['state_classes']

print(f"Spectrograms shape: {X_memmap.shape}")
print(f"Labels shape: {y_drone.shape}")
print(f"File IDs shape: {file_ids.shape} (unique files: {len(np.unique(file_ids))})")
print(f"Drone classes: {drone_classes}")
print(f"Interference classes: {interference_classes}")
print(f"State classes: {state_classes}")
print(f"Number of classes: {len(drone_classes)}")
```

Spectrograms shape: (19478, 224, 224, 3)
 Labels shape: (19478,)
 File IDs shape: (19478,) (unique files: 195)
 Drone classes: ['AIR' 'DIS' 'INS' 'MA1' 'MAV' 'MIN' 'PHA']
 Interference classes: ['BOTH' 'CLEAN']
 State classes: ['FY' 'HO' 'ON']
 Number of classes: 7

7. Train/Test Split

```
In [30]: print("Performing file-level stratified split...")
train_idx, test_idx = get_stratified_file_split(
    X_memmap, y_drone, file_ids,
    test_size=CONFIG['test_size'],
    random_state=CONFIG['random_state']
)

# Split labels only (X handled by RGBMemmapDataset)
y_train = y_drone[train_idx]
y_test = y_drone[test_idx]
y_interference_test = y_interference[test_idx]
y_state_test = y_state[test_idx]

print(f"Train set: {len(train_idx)} samples")
print(f"Test set: {len(test_idx)} samples")
print(f"Train class distribution: {np.bincount(y_train)}")
print(f"Test class distribution: {np.bincount(y_test)}")
print(f"\nMemory saved: Indices stored instead of 11.73 GB array copy")

# Save test data for reuse
test_data_dir = CONFIG['test_data_dir']
os.makedirs(test_data_dir, exist_ok=True)

# Load full test set into memory from memmap
print("\nLoading test set into memory for saving...")
X_test = X_memmap[test_idx].copy()

# Save full test data
test_data_path = os.path.join(test_data_dir, 'cnn_test_data.npz')
np.savez(
    test_data_path,
    X_test=X_test,
    y_test=y_test,
    y_interference_test=y_interference_test,
    y_state_test=y_state_test,
    test_idx=test_idx,
    file_ids_test=file_ids[test_idx],
    drone_classes=drone_classes,
    interference_classes=interference_classes,
```

```

        state_classes=state_classes
    )
print(f"Full test data saved to {test_data_path}")

# Save separated files per Drone and Interference (Hierarchical)
print("\nGenerating separated test files (structure: spectrogram/INT/DRONE/)...")

for d_idx, drone_class in enumerate(drone_classes):
    for i_idx, int_class in enumerate(interference_classes):
        # Filter for specific drone and interference
        mask = (y_test == d_idx) & (y_interference_test == i_idx)

        if not np.any(mask):
            continue

        X_sub = X_test[mask]
        y_sub = y_test[mask]
        y_int_sub = y_interference_test[mask]

        # Define components for hierarchy and filename
        data_type = 'spectrogram'
        int_name = str(int_class)
        drone_name = str(drone_class)
        dims = "224x224x3"

        # Create directory structure: output/test_data/{INT}/
        save_dir = os.path.join(test_data_dir, int_name)
        os.makedirs(save_dir, exist_ok=True)

        # Construct filename: spectrogram_{INT}_{DRONE}_224x224x3.npz
        filename = f"{data_type}_{int_name}_{drone_name}_{dims}.npz"
        file_path = os.path.join(save_dir, filename)

        np.savez(
            file_path,
            X=X_sub,
            y=y_sub,
            y_interference=y_int_sub,
            drone_class=drone_class,
            interference_class=int_class

```

```
)  
    print(f"  Saved {filename} in {save_dir} ({len(X_sub)} samples)")
```

Performing file-level stratified split...
Train set: 15587 samples
Test set: 3891 samples
Train class distribution: [2400 1600 2387 2400 2300 2400 2100]
Test class distribution: [600 400 591 600 700 600 400]

Memory saved: Indices stored instead of 11.73 GB array copy

Loading test set into memory for saving...
Full test data saved to drive/MyDrive/DroneDetect_V2/output/sample/test_data/cnn_test_data.npz

Generating separated test files (structure: spectrogram/INT/DRONE/)...

Saved spectrogram_BOTH_AIR_224x224x3.npz in drive/MyDrive/DroneDetect_V2/output/sample/test_data/BOTH (400 samples)
Saved spectrogram_CLEAN_AIR_224x224x3.npz in drive/MyDrive/DroneDetect_V2/output/sample/test_data/CLEAN (200 samples)
Saved spectrogram_BOTH_DIS_224x224x3.npz in drive/MyDrive/DroneDetect_V2/output/sample/test_data/BOTH (100 samples)
Saved spectrogram_CLEAN_DIS_224x224x3.npz in drive/MyDrive/DroneDetect_V2/output/sample/test_data/CLEAN (300 samples)
Saved spectrogram_BOTH_INS_224x224x3.npz in drive/MyDrive/DroneDetect_V2/output/sample/test_data/BOTH (291 samples)
Saved spectrogram_CLEAN_INS_224x224x3.npz in drive/MyDrive/DroneDetect_V2/output/sample/test_data/CLEAN (300 samples)
Saved spectrogram_BOTH_MA1_224x224x3.npz in drive/MyDrive/DroneDetect_V2/output/sample/test_data/BOTH (300 samples)
Saved spectrogram_CLEAN_MA1_224x224x3.npz in drive/MyDrive/DroneDetect_V2/output/sample/test_data/CLEAN (300 samples)
Saved spectrogram_BOTH_MAV_224x224x3.npz in drive/MyDrive/DroneDetect_V2/output/sample/test_data/BOTH (400 samples)
Saved spectrogram_CLEAN_MAV_224x224x3.npz in drive/MyDrive/DroneDetect_V2/output/sample/test_data/CLEAN (300 samples)
Saved spectrogram_BOTH_MIN_224x224x3.npz in drive/MyDrive/DroneDetect_V2/output/sample/test_data/BOTH (300 samples)
Saved spectrogram_CLEAN_MIN_224x224x3.npz in drive/MyDrive/DroneDetect_V2/output/sample/test_data/CLEAN (300 samples)
Saved spectrogram_BOTH_PHA_224x224x3.npz in drive/MyDrive/DroneDetect_V2/output/sample/test_data/BOTH (200 samples)
Saved spectrogram_CLEAN_PHA_224x224x3.npz in drive/MyDrive/DroneDetect_V2/output/sample/test_data/CLEAN (200 samples)

8. Create DataLoaders

```
In [31]: train_dataset = RGBMemmapDataset(X_memmap, train_idx, y_train)  
test_dataset = RGBMemmapDataset(X_memmap, test_idx, y_test)  
  
train_loader = DataLoader(  
    train_dataset,  
    batch_size=CONFIG['batch_size'],  
    shuffle=True,  
    num_workers=2,
```

```

        pin_memory=True
    )

test_loader = DataLoader(
    test_dataset,
    batch_size=CONFIG['batch_size'],
    shuffle=False,
    num_workers=2,
    pin_memory=True
)

print(f"Train batches: {len(train_loader)}")
print(f"Test batches: {len(test_loader)}")
print(f"Memory per batch: ~{CONFIG['batch_size'] * 224 * 224 * 3 * 4 / 1024**2:.1f} MB")

```

Train batches: 122
 Test batches: 31
 Memory per batch: ~73.5 MB

8b. Memory Profiling Checkpoint

```

In [32]: process = psutil.Process()
mem_info = process.memory_info()
print(f"Memory usage: {mem_info.rss / 1024**3:.2f} GB")

# Force garbage collection
gc.collect()
if torch.cuda.is_available():
    torch.cuda.empty_cache()
    print(f"GPU memory allocated: {torch.cuda.memory_allocated() / 1024**3:.2f} GB")
    print(f"GPU memory reserved: {torch.cuda.memory_reserved() / 1024**3:.2f} GB")

```

Memory usage: 15.50 GB
 GPU memory allocated: 0.20 GB
 GPU memory reserved: 0.22 GB

9. Training Function

```
In [33]: def train_model(model, train_loader, test_loader, epochs=10, lr=0.01, device='cuda'):  
    """  
        Train a PyTorch model with memory cleanup.  
  
    Parameters  
    -----  
    model : nn.Module  
        Model to train  
    train_loader : DataLoader  
        Training data loader  
    test_loader : DataLoader  
        Test data loader  
    epochs : int  
        Number of epochs  
    lr : float  
        Learning rate  
    device : str  
        Device to use ('cuda' or 'cpu')  
  
    Returns  
    -----  
    model : nn.Module  
        Trained model  
    history : dict  
        Training history  
    """  
    model = model.to(device)  
    criterion = nn.CrossEntropyLoss()  
    optimizer = optim.Adam(model.parameters(), lr=lr)  
  
    history = {  
        'train_loss': [],  
        'train_acc': [],  
        'test_loss': [],  
        'test_acc': []  
    }  
  
    for epoch in range(epochs):  
        # Training phase  
        model.train()
```

```
train_loss = 0.0
train_correct = 0
train_total = 0

for batch_X, batch_y in tqdm(train_loader, desc=f"Epoch {epoch+1}/{epochs}"):
    batch_X = batch_X.to(device)
    batch_y = batch_y.to(device)

    optimizer.zero_grad()
    outputs = model(batch_X)
    loss = criterion(outputs, batch_y)
    loss.backward()
    optimizer.step()

    train_loss += loss.item()
    _, predicted = torch.max(outputs, 1)
    train_total += batch_y.size(0)
    train_correct += (predicted == batch_y).sum().item()

# Validation phase
model.eval()
test_loss = 0.0
test_correct = 0
test_total = 0

with torch.no_grad():
    for batch_X, batch_y in test_loader:
        batch_X = batch_X.to(device)
        batch_y = batch_y.to(device)

        outputs = model(batch_X)
        loss = criterion(outputs, batch_y)

        test_loss += loss.item()
        _, predicted = torch.max(outputs, 1)
        test_total += batch_y.size(0)
        test_correct += (predicted == batch_y).sum().item()

# Record history
history['train_loss'].append(train_loss / len(train_loader))
history['train_acc'].append(100 * train_correct / train_total)
```

```

        history['test_loss'].append(test_loss / len(test_loader))
        history['test_acc'].append(100 * test_correct / test_total)

        print(f"Epoch {epoch+1}/{epochs}: "
              f"Train Loss: {history['train_loss'][-1]:.4f}, "
              f"Train Acc: {history['train_acc'][-1]:.2f}%, "
              f"Test Loss: {history['test_loss'][-1]:.4f}, "
              f"Test Acc: {history['test_acc'][-1]:.2f}%)"

    # Memory cleanup after each epoch
    gc.collect()
    if torch.cuda.is_available():
        torch.cuda.empty_cache()

return model, history

print("Training function defined")

```

Training function defined

10. Train VGG16

```

In [34]: num_classes = len(drone_classes)
vgg_model = VGG16FC(num_classes=num_classes)

print(f"Training VGG16 with {num_classes} classes...")
vgg_model, vgg_history = train_model(
    vgg_model,
    train_loader,
    test_loader,
    epochs=CONFIG['epochs'],
    lr=CONFIG['learning_rate'],
    device=CONFIG['device']
)
print("VGG16 training complete!")

```

Training VGG16 with 7 classes...

Epoch 1/10: 100%|██████████| 122/122 [00:08<00:00, 14.20it/s]

Epoch 1/10: Train Loss: 1.8804, Train Acc: 69.26%, Test Loss: 1.1507, Test Acc: 78.98%

```
Epoch 2/10: 100%|██████████| 122/122 [00:08<00:00, 14.17it/s]
Epoch 2/10: Train Loss: 0.6801, Train Acc: 82.84%, Test Loss: 1.0365, Test Acc: 80.83%
Epoch 3/10: 100%|██████████| 122/122 [00:08<00:00, 14.16it/s]
Epoch 3/10: Train Loss: 0.4813, Train Acc: 86.71%, Test Loss: 1.5858, Test Acc: 75.35%
Epoch 4/10: 100%|██████████| 122/122 [00:08<00:00, 14.13it/s]
Epoch 4/10: Train Loss: 0.3898, Train Acc: 88.99%, Test Loss: 1.3841, Test Acc: 80.98%
Epoch 5/10: 100%|██████████| 122/122 [00:08<00:00, 14.22it/s]
Epoch 5/10: Train Loss: 0.2737, Train Acc: 91.71%, Test Loss: 1.6196, Test Acc: 79.36%
Epoch 6/10: 100%|██████████| 122/122 [00:08<00:00, 14.23it/s]
Epoch 6/10: Train Loss: 0.2507, Train Acc: 92.40%, Test Loss: 1.9620, Test Acc: 79.21%
Epoch 7/10: 100%|██████████| 122/122 [00:08<00:00, 14.21it/s]
Epoch 7/10: Train Loss: 0.3509, Train Acc: 91.26%, Test Loss: 1.8645, Test Acc: 79.75%
Epoch 8/10: 100%|██████████| 122/122 [00:08<00:00, 14.22it/s]
Epoch 8/10: Train Loss: 0.1857, Train Acc: 94.48%, Test Loss: 2.0230, Test Acc: 76.87%
Epoch 9/10: 100%|██████████| 122/122 [00:08<00:00, 14.23it/s]
Epoch 9/10: Train Loss: 0.1481, Train Acc: 95.42%, Test Loss: 1.9568, Test Acc: 78.87%
Epoch 10/10: 100%|██████████| 122/122 [00:08<00:00, 14.20it/s]
Epoch 10/10: Train Loss: 0.1645, Train Acc: 95.11%, Test Loss: 2.0290, Test Acc: 78.51%
VGG16 training complete!
```

11. VGG16 Performance Metrics

```
In [35]: vgg_model.eval()
vgg_preds = []
vgg_true = []

with torch.no_grad():
    for batch_X, batch_y in test_loader:
        batch_X = batch_X.to(CONFIG['device'])
        outputs = vgg_model(batch_X)
        _, predicted = torch.max(outputs, 1)
        vgg_preds.extend(predicted.cpu().numpy())
        vgg_true.extend(batch_y.numpy())

vgg_preds = np.array(vgg_preds)
vgg_true = np.array(vgg_true)
```

```

print("VGG16 Classification Report:")
print(classification_report(vgg_true, vgg_preds, target_names=drone_classes))
print(f"\nVGG16 Accuracy: {accuracy_score(vgg_true, vgg_preds):.4f}")
print(f"VGG16 F1 Score (macro): {f1_score(vgg_true, vgg_preds, average='macro'): .4f}")

```

VGG16 Classification Report:

	precision	recall	f1-score	support
AIR	0.71	0.76	0.73	600
DIS	0.97	0.86	0.91	400
INS	0.87	0.75	0.80	591
MA1	0.65	0.72	0.69	600
MAV	0.67	0.56	0.61	700
MIN	1.00	1.00	1.00	600
PHA	0.73	0.97	0.83	400
accuracy			0.79	3891
macro avg	0.80	0.80	0.80	3891
weighted avg	0.79	0.79	0.78	3891

VGG16 Accuracy: 0.7851

VGG16 F1 Score (macro): 0.7962

12. VGG16 Confusion Matrix

```

In [36]: cm_vgg = confusion_matrix(vgg_true, vgg_preds)

# Create confusion matrix heatmap with plotly
fig = go.Figure(data=go.Heatmap(
    z=cm_vgg,
    x=list(drone_classes),
    y=list(drone_classes),
    colorscale='Blues',
    text=cm_vgg,
    texttemplate='%{text}',
    textfont={'size': 12},
    hoverongaps=False
))

```

```
fig.update_layout(  
    title='VGG16 Confusion Matrix',  
    xaxis_title='Predicted Label',  
    yaxis_title='True Label',  
    xaxis={'side': 'bottom'},  
    yaxis={'autorange': 'reversed'},  
    width=800,  
    height=700  
)  
fig.show()  
save_figure(fig)
```


Saved: figures/training_cnn_COLAB/VGG16_Confusion_Matrix.png

13. Train ResNet50

```
In [37]: num_classes = len(drone_classes)
resnet_model = ResNet50FC(num_classes=num_classes)

print(f"Training ResNet50 with {num_classes} classes...")
resnet_model, resnet_history = train_model(
    resnet_model,
    train_loader,
    test_loader,
    epochs=CONFIG['epochs'],
    lr=CONFIG['learning_rate'],
    device=CONFIG['device']
)
print("ResNet50 training complete!")
```

Training ResNet50 with 7 classes...

```
Epoch 1/10: 100%|██████████| 122/122 [00:05<00:00, 20.55it/s]
Epoch 1/10: Train Loss: 70.6483, Train Acc: 59.32%, Test Loss: 13.0459, Test Acc: 74.15%
Epoch 2/10: 100%|██████████| 122/122 [00:05<00:00, 20.65it/s]
Epoch 2/10: Train Loss: 12.1760, Train Acc: 74.65%, Test Loss: 12.9827, Test Acc: 73.22%
Epoch 3/10: 100%|██████████| 122/122 [00:05<00:00, 20.63it/s]
Epoch 3/10: Train Loss: 10.6961, Train Acc: 79.26%, Test Loss: 21.0434, Test Acc: 70.26%
Epoch 4/10: 100%|██████████| 122/122 [00:05<00:00, 20.63it/s]
Epoch 4/10: Train Loss: 7.6036, Train Acc: 83.60%, Test Loss: 27.3043, Test Acc: 71.16%
Epoch 5/10: 100%|██████████| 122/122 [00:05<00:00, 20.56it/s]
Epoch 5/10: Train Loss: 8.6377, Train Acc: 83.79%, Test Loss: 19.4189, Test Acc: 73.81%
Epoch 6/10: 100%|██████████| 122/122 [00:05<00:00, 20.62it/s]
Epoch 6/10: Train Loss: 4.2715, Train Acc: 89.77%, Test Loss: 27.0581, Test Acc: 72.94%
Epoch 7/10: 100%|██████████| 122/122 [00:05<00:00, 20.64it/s]
Epoch 7/10: Train Loss: 5.4503, Train Acc: 88.74%, Test Loss: 25.0186, Test Acc: 72.73%
Epoch 8/10: 100%|██████████| 122/122 [00:05<00:00, 20.60it/s]
Epoch 8/10: Train Loss: 4.9373, Train Acc: 89.34%, Test Loss: 22.7478, Test Acc: 76.23%
Epoch 9/10: 100%|██████████| 122/122 [00:05<00:00, 20.64it/s]
```

Epoch 9/10: Train Loss: 3.3305, Train Acc: 92.24%, Test Loss: 30.9102, Test Acc: 69.73%

Epoch 10/10: 100%|██████████| 122/122 [00:05<00:00, 20.62it/s]

Epoch 10/10: Train Loss: 5.6662, Train Acc: 89.86%, Test Loss: 30.9221, Test Acc: 73.68%

ResNet50 training complete!

14. ResNet50 Performance Metrics

```
In [38]: resnet_model.eval()
resnet_preds = []
resnet_true = []

with torch.no_grad():
    for batch_X, batch_y in test_loader:
        batch_X = batch_X.to(CONFIG['device'])
        outputs = resnet_model(batch_X)
        _, predicted = torch.max(outputs, 1)
        resnet_preds.extend(predicted.cpu().numpy())
        resnet_true.extend(batch_y.numpy())

resnet_preds = np.array(resnet_preds)
resnet_true = np.array(resnet_true)

print("ResNet50 Classification Report:")
print(classification_report(resnet_true, resnet_preds, target_names=drone_classes))
print(f"\nResNet50 Accuracy: {accuracy_score(resnet_true, resnet_preds):.4f}")
print(f"ResNet50 F1 Score (macro): {f1_score(resnet_true, resnet_preds, average='macro'): .4f}")
```

ResNet50 Classification Report:

	precision	recall	f1-score	support
AIR	0.82	0.72	0.77	600
DIS	0.73	0.94	0.82	400
INS	0.70	0.80	0.75	591
MA1	0.54	0.74	0.62	600
MAV	0.64	0.40	0.49	700
MIN	1.00	0.97	0.99	600
PHA	0.87	0.69	0.77	400
accuracy			0.74	3891
macro avg	0.76	0.75	0.74	3891
weighted avg	0.75	0.74	0.73	3891

ResNet50 Accuracy: 0.7368

ResNet50 F1 Score (macro): 0.7433

15. ResNet50 Confusion Matrix

```
In [39]: cm_resnet = confusion_matrix(resnet_true, resnet_preds)

# Create confusion matrix heatmap with plotly
fig = go.Figure(data=go.Heatmap(
    z=cm_resnet,
    x=list(drone_classes),
    y=list(drone_classes),
    colorscale='Greens',
    text=cm_resnet,
    texttemplate=' %{text} ',
    textfont={'size': 12},
    hoverongaps=False
))

fig.update_layout(
    title='ResNet50 Confusion Matrix',
    xaxis_title='Predicted Label',
    yaxis_title='True Label',
```

```
xaxis={'side': 'bottom'},  
yaxis={'autorange': 'reversed'},  
width=800,  
height=700  
)  
fig.show()  
save_figure(fig)
```


Saved: figures/training_cnn_COLAB/ResNet50_Confusion_Matrix.png

16. Side-by-Side Performance Comparison

```
In [40]: # Training curves comparison
fig = make_subplots(rows=1, cols=2, subplot_titles=('Training History Comparison', 'Final Performance Comparison'))

# Training curves
epochs = list(range(1, len(vgg_history['train_acc']) + 1))

fig.add_trace(go.Scatter(x=epochs, y=vgg_history['train_acc'], mode='lines+markers', name='VGG16 Train', line=dict(color='blue'))
fig.add_trace(go.Scatter(x=epochs, y=vgg_history['test_acc'], mode='lines+markers', name='VGG16 Test', line=dict(color='blue'),
fig.add_trace(go.Scatter(x=epochs, y=resnet_history['train_acc'], mode='lines+markers', name='ResNet50 Train', line=dict(color='red'))
fig.add_trace(go.Scatter(x=epochs, y=resnet_history['test_acc'], mode='lines+markers', name='ResNet50 Test', line=dict(color='red'))

# Final metrics comparison
models = ['VGG16', 'ResNet50']
accuracies = [accuracy_score(vgg_true, vgg_preds), accuracy_score(resnet_true, resnet_preds)]
f1_scores = [f1_score(vgg_true, vgg_preds, average='macro'), f1_score(resnet_true, resnet_preds, average='macro')]

fig.add_trace(go.Bar(x=models, y=accuracies, name='Accuracy', marker_color='steelblue'), row=1, col=2)
fig.add_trace(go.Bar(x=models, y=f1_scores, name='F1 Score (macro)', marker_color='coral'), row=1, col=2)

fig.update_layout(
    title='CNN Training Comparison - VGG16 vs ResNet50',
    height=500,
    width=1200,
    barmode='group'
)
fig.update_xaxes(title_text='Epoch', row=1, col=1)
fig.update_yaxes(title_text='Accuracy (%)', row=1, col=1)
fig.update_xaxes(title_text='Model', row=1, col=2)
fig.update_yaxes(title_text='Score', row=1, col=2)

fig.show()
save_figure(fig)

print("\nFinal Results:")
```

```
print(f"VGG16 - Accuracy: {accuracies[0]:.4f}, F1: {f1_scores[0]:.4f}")
print(f"ResNet50 - Accuracy: {accuracies[1]:.4f}, F1: {f1_scores[1]:.4f}")
```



Saved: figures/training_cnn_COLAB/CNN_Training_Comparison_VGG16_vs_ResNet50.png



Final Results:

VGG16 - Accuracy: 0.7851, F1: 0.7962
ResNet50 - Accuracy: 0.7368, F1: 0.7433

17. Save Models

```
In [41]: os.makedirs(CONFIG['models_dir'], exist_ok=True)

# Save VGG16
vgg_path = os.path.join(CONFIG['models_dir'], 'vgg16_cnn.pth')
torch.save({
    'model_state_dict': vgg_model.state_dict(),
    'history': vgg_history,
    'num_classes': num_classes,
    'drone_classes': drone_classes
}, vgg_path)
print(f"VGG16 model saved to {vgg_path}")

# Save ResNet50
resnet_path = os.path.join(CONFIG['models_dir'], 'resnet50_cnn.pth')
torch.save({
    'model_state_dict': resnet_model.state_dict(),
    'history': resnet_history,
    'num_classes': num_classes,
    'drone_classes': drone_classes
}, resnet_path)
print(f"ResNet50 model saved to {resnet_path}")
```

```
VGG16 model saved to drive/MyDrive/DroneDetect_V2/output/models/vgg16_cnn.pth
ResNet50 model saved to drive/MyDrive/DroneDetect_V2/output/models/resnet50_cnn.pth
```

18. Summary

```
In [42]: print("=" * 60)
print("CNN TRAINING SUMMARY")
print("=" * 60)

print(f"\nDataset: {len(y_drone)} total | {len(y_train)} train | {len(y_test)} test | {num_classes} classes")
print(f"Training: {CONFIG['epochs']} epochs, batch_size={CONFIG['batch_size']}, lr={CONFIG['learning_rate']}, device={CONFIG['device']}")

print(f"\nVGG16: Accuracy={accuracy_score(vgg_true, vgg_preds):.4f}, F1={f1_score(vgg_true, vgg_preds, average='macro'):.4f}")
```

```
print(f"ResNet50: Accuracy={accuracy_score(resnet_true, resnet_preds):.4f}, F1={f1_score(resnet_true, resnet_preds, average='micro'): .4f} ")
print(f"\nModels saved to: {CONFIG['models_dir']}")
print("=" * 60)
```

```
=====
```

```
CNN TRAINING SUMMARY
```

```
=====
```

```
Dataset: 19478 total | 15587 train | 3891 test | 7 classes
```

```
Training: 10 epochs, batch_size=128, lr=0.01, device=cuda
```

```
VGG16: Accuracy=0.7851, F1=0.7962
```

```
ResNet50: Accuracy=0.7368, F1=0.7433
```

```
Models saved to: drive/MyDrive/DroneDetect_V2/output/models/
=====
```